

PersonaPage

Product Foundation Document

Phase 1 — Product Definition · Working title: IdentityOS

"A world where every link you share shows exactly the right version of you to exactly the right person."

Key Decisions

Primary Outcome	Stronger personal brand
Wedge User	CS students & early-career developers
Core Mechanic	Context-first smart links
Personal Pain	Same CV/portfolio sent to everyone
Positioning	Adaptive identity platform

Contents

01 Vision Statement	06 Competitive Analysis
02 Problem Statement	07 MVP Scope
03 Ideal Customer Profile	08 Feature Prioritization
04 User Personas	09 Non-Goals
05 Core Product Thesis	10 90-Day Roadmap

Vision Statement

One-liner

"A world where every link you share shows exactly the right version of you to exactly the right person."

Expanded

Professionals lose opportunities every day not because they lack skills — but because they present the same static identity to every audience. PersonaPage makes your professional identity as contextual as a real conversation: the person opening your link gets the version of you most relevant to them, automatically.

The contrarian bet

"The problem isn't that people can't write about themselves. The problem is that one link was never supposed to work for everyone. We believe professional identity should be audience-aware by default."

Problem Statement

The Situation	CS students, early-career developers, and builders maintain one version of their portfolio/resume and share it universally — to startup founders, FAANG recruiters, design leads, and collaborators alike.
The Trigger	The exact moment you copy-paste your GitHub or portfolio link into a cold message, job application, or hackathon bio and think: "This doesn't quite fit what they're looking for."
Current Behaviour	People either: (a) maintain 3–4 manually updated versions, (b) send the generic version hoping it lands, or (c) spend 30+ minutes rewriting their bio for each outreach.
Why Tools Fail	LinkedIn: one public profile, no adaptation. Notion: static, one view. Resume builders: file output not a live link. Read.cv / Polywork: better aesthetics, still single-version. None treat the viewer as a variable.
The Cost	Tangible: missed interviews, poor cold-outreach response rates. Intangible: the quiet anxiety of knowing your work is good but your presentation doesn't do it justice.
Frequency	Every job application. Every cold LinkedIn message. Every hackathon intro. For an active networker: 5–15 times per week.

Ideal Customer Profile

Primary Segment: CS students in their final 1–2 years, or recent grads (0–2 years experience), actively job hunting, freelancing, or building side projects — who care about being taken seriously as builders.

Demographics & psychographics

Age range	20–26
Career stage	Student → junior / early career
Technical level	Comfortable with tools; wants zero-config UX
Context	University, hackathons, tech Twitter/X, Discord builder communities
Cares about	Being perceived as a real builder, not just a student
Afraid of	Being overlooked because their presentation is weak
Believes	Their actual work is better than their current CV shows

Behavioural signals

- Has a GitHub but no clean portfolio page
- Has applied to 10+ jobs with the same CV
- Actively attends hackathons or tech meetups
- Follows build-in-public culture on Twitter/X
- Has tried Notion, Linktree, or Read.cv at some point

Adoption triggers vs. barriers

Will switch when	Gets ghosted after sending a generic GitHub link and realises it told the wrong story to the wrong person.
Won't adopt if	AI-generated content doesn't sound like them, or setup takes more than 10 minutes.
Where to find them	Uni CS Discord/Slack, hackathon lists (MLH, Devpost), r/cscareerquestions, #buildinpublic on X.

User Personas

Aryan — The Overlooked Builder

Final-year CS student, building side projects, applying to startups

"His GitHub has 40 repos. His CV makes him look like every other student."

Primary goal	Land a SWE internship or junior role at a startup
Secondary goal	Be taken seriously at hackathons and builder events
Frustration 1	Startup recruiters want projects; FAANG wants structured CV — same link fails both
Frustration 2	No way to know if anyone actually looked at his portfolio link
Sign-up trigger	Gets rejected after a strong technical round. Realises his portfolio led with coursework, not his best project.
Success	Gets a reply from cold outreach within 48 hours — his AI project link sent to an AI startup recruiter actually landed.
Churn risk	If AI-generated versions sound generic and don't reflect his actual voice.

Priya — The Multi-Hat Founder

Recent CS grad, freelancing as a dev, building a side startup, looking for a co-founder

"She's three different people professionally — and has one LinkedIn that captures none of them."

Primary goal	Find a technical co-founder who matches her vision
Secondary goal	Land 2–3 freelance clients to fund runway
Frustration 1	Founder networking requires a completely different story than client pitching
Frustration 2	Manually maintaining multiple versions is exhausting with unclear ROI
Sign-up trigger	Heading to a tech meetup. Needs to share her profile with a potential client AND co-founder in the same evening.
Success	Creates two smart links before the event. Follows up knowing exactly which version each person saw.
Churn risk	If creating and managing multiple versions becomes more work than writing manually.

Core Product Thesis

The Contrarian Belief

Every other portfolio tool optimises for how YOU present yourself. PersonaPage optimises for how THEY perceive you. The viewer is the variable — not the content.

The Insight That Makes It Now (2025)

LLMs can model audience intent with enough fidelity to reshape narrative framing without losing authentic voice. Combined with link-level routing and lightweight profiling, you can build audience-aware identity without manual work per viewer.

The Core Mechanic

User creates ONE master profile (projects, skills, experience, tone). PersonaPage generates N smart links — each pre-configured for a different audience context. Each link renders a different prioritisation, framing, and narrative from the same underlying data.

The Compounding Mechanic

More projects added → richer source material per version. More links shared → analytics data on what resonates per context. The product gets more valuable the more someone uses it.

The 90-Day Falsifiability Test

If users create an account but never share a second contextual link beyond their first — the thesis is wrong. It means people don't need multiple versions; they just need ONE better version. This would pivot you toward a smarter resume builder, not an adaptive identity platform.

Competitive Analysis

Tool	Does well	Misses
LinkedIn	Distribution & credibility	One profile, zero adaptation
Read.cv	Clean aesthetic, builder community	Static, no AI layer
Notion	Flexible layout, free tier	No audience awareness
Polywork	Multi-role identity framing	Manual, no smart links
Resume builders	CV formatting & templates	File output, not a live link
Linktree	Simple link aggregation	Zero identity depth
Gamma / Tome	AI-powered presentations	Not identity-focused

Your gap: Nobody owns 'the smart link that shows the right version of you.' Linktree owns links. LinkedIn owns identity. Nobody owns the intersection: live link + adaptive profile + audience context + analytics.

Honest threat: LinkedIn could ship 'audience modes' tomorrow. Your moat in Year 1 is NOT technology — it's UX focus and speed of iteration for this specific user pain.

0
7

MVP Scope

The core loop — the only thing that matters

0
1**User inputs their background**

Projects, skills, experience, tone — structured form

0
2**Define 2–3 audience contexts**

e.g. startup recruiter, technical co-founder, freelance client

0
3**AI generates tailored views**

Different prioritisation, framing, and narrative per context

0
4**Each view gets a unique link**

`yourname.personapage.io/for/recruiters`

0
5**User shares the right link**

Right link to right person at the right moment

0
6**Basic analytics**

Who opened it, how long, from where

Technical requirements

- ✓ Auth + user profile (Next.js + Supabase)
- ✓ Master profile input form (projects, skills, bio, tone — structured fields)
- ✓ Audience context selector (predefined templates to start)
- ✓ LLM integration to generate contextual framing (Claude / GPT-4o)
- ✓ Dynamic page renderer per link / context
- ✓ Unique URL per context — `yourname.personapage.io/for/[context]`
- ✓ Basic link analytics: view count, timestamp, referrer

Does NOT require for MVP: Real-time audience auto-detection · Custom domains · Mobile app · Team/org features · GitHub sync · Payments · A/B testing of versions

0
8

Feature Prioritization

TIER 1 — Build first (Week 1–4)

- ✓ Master profile builder (text input, structured fields)
- ✓ Audience context templates (3 defaults: recruiter, founder, client)
- ✓ AI profile renderer per context
- ✓ Unique shareable link per context
- ✓ Basic view analytics (opens, time on page)

TIER 2 — After first 10 users give feedback (Week 5–8)

- ✓ Custom context creator (user defines their own audience type)
- ✓ GitHub project import (auto-pull repo descriptions)
- ✓ Visual profile customisation (colours, layout)
- ✓ Link expiry / access control

TIER 3 — Only if thesis is validated (Month 3+)

- ✓ Viewer intent detection (clicks, time per section)
- ✓ AI tone calibration (formal vs casual per context)
- ✓ LinkedIn data import
- ✓ Embeddable widget for existing portfolios

NEVER BUILD (MVP)

- ✗ Mobile app
- ✗ Team / org accounts
- ✗ Real-time collaboration
- ✗ Marketplace / directory of profiles

09

Non-Goals

Not a resume formatter	Other tools do this. PersonaPage is a live link, not a file.
Not a job board	No recruiter-facing discovery or job listings.
Not a social network	No feed, following, or profile discovery in v1.
Not a LinkedIn replacement	We work alongside LinkedIn, not against it.
Not for enterprises	Solo professionals only for MVP. No org accounts.
Not for non-technical users	v1 is built for and validated with technical builders.
No auto-viewer detection	Viewer intent detection is a research problem. Let users pre-configure contexts — that is the product problem.

10

90-Day Roadmap

Month 1	Month 2	Month 3
Build	Ship + Learn	Validate or Pivot
• Week 1–2: Core profile input + Supabase schema	• Week 5: Launch to 10 people personally	• Watch: do users share 3+ links? (falsifiability test)
• Week 3: LLM integration + context rendering engine	• Week 6: Watch usage — what do they share? what breaks?	• → YES: thesis confirmed — add analytics depth
• Week 4: Shareable links + basic analytics dashboard	• Week 7: Fix the top 3 friction points • Week 8: Add GitHub import + custom context creator	• → NO: pivot to best-single AI portfolio for CS students • Either way: document learnings as a public case study

Portfolio & career goals by end of Month 3

- ✓ A live product with real users — not a tutorial clone
- ✓ A case study: "I identified this problem, built this, learned this, changed this"
- ✓ A story for internship interviews that positions you as an AI product engineer
- ✓ A demonstrable thesis: different links → different versions → different outcomes

The one thing to internalise: You don't need a perfect product. You need a demonstrable thesis. PersonaPage's MVP must answer one question in 30 seconds: "You send different links to different people, and each link shows a different version of you." If anyone understands that immediately — you have a product. Everything else is polish.